

Lec01: Stack Overflow and Protections

Taesoo Kim

Goals and Lessons

- Learn about the stack overflow bugs
- Understand their security implications (i.e., control hijacking)
- Understand the off-the-shelf mitigation (i.e., ssp, DEP, RELO, ASLR)
- Learn them from the real-world examples (i.e., qemu/ruby/wireshark)

“Smashing The Stack For Fun And Profit”

CS101: What's Wrong?

```
main() {  
    char buf[16];  
    scanf("%s", buf);  
}
```

CS101: What's Wrong?

```
main() {  
    char buf[16];  
    scanf("%s", buf);  
}
```

1. `scanf("%s", &buf)`
2. `scanf("%16s", buf)`
3. `scanf("%15s", buf)`
4. `scanf("%as", &bufptr)`

Error-prone C APIs: scanf()

- `scanf("%15s", buf) // CORRECT!`

```
$ cd lec01-stackovfl/apis
$ cat scanf.c
$ make
$ ./scanf
...
```

Error-prone C APIs: scanf()

- `scanf("%16s", buf) // BUG!`

```
$ ./scanf
aaaaaaaaaaaaaaaaaaaaaaaaaaaaa
aaaaaaaaaaaaaaaaaaaaa
01: 61 (a)
02: 61 (a)
...
0e: 61 (a)
0f: 61 (a)
10: 61 (a)
11: 00 ()
12: BE ()
13: AD ()
14: DE ()
```

Control Hijacking Attack: Basic Idea

```
main() {  
    char buf[16];  
    scanf("%s", buf); // BUG!  
}
```

```
(top, growing)  
|<--- current frame -->|<-- caller-->|  
[buf          ][fp][ra][...          ]  
|<----- 16 ---->|
```


Control Hijacking Attack: Basic Idea

```
main() {  
    char buf[16];  
    scanf("%s", buf); // BUG!  
}
```

```
(top, growing)  
|<--- current frame -->|<-- caller-->|  
[buf                ][fp][ra][...      ]  
|<---- 16 ---->|  
AAAAAAAAAAAAAAAAABBBBCCCC... =>
```

```
!SEGFALT @eip=CCCC
```

Control Hijacking Attack: Where to Jump?

- Jump to the injected code (e.g., stack, environ, etc)

```
(top, growing)
|<--- current frame -->|<-- caller-->|
[buf                ][fp][ra][...      ]
|<----- 16 ---->|
AAAAAAAAAAAAAAAAABBBBXXX[shellcode ..]
                    |      ^
                    +-----+ (1) attacker's input
```

Control Hijacking Attack: Where to Jump?

- Jump to the injected code (e.g., stack, environ, etc)

```
(top, growing)
|<--- current frame -->|
[buf                ][fp][ra] ... [SHELLCODE=...]
|<----- 16 ---->|          ^
AAAAAAAAAAAAAAAAABBBBXXXX    |
                        |      |
                        +-----+ (2) crafted environ
```

Control Hijacking Attack: Advanced Topics

- Stack pivoting when frame pointer is crafted
- Even off-by-one (e.g., writing NULL) is enough for stack pivoting

(top, growing)

|<--- current frame -->|<-- caller-->|

[buf [fp][ra][...]

|<--- 16 --->|

(1)AAAAAAAAAAAAAAAAAXXX -> enough to control (i.e., leave; *ret*)

(2)AAAAAAAAAAAAAAAAAX -> off-by-one

Memory Safety in C/C++

- Spatial safety → e.g., buffer over/underflow
- Temporal safety → e.g., use-after-free

Addressing Memory Safety Issues in C/C++

- Spatial safety → e.g., buffer over/underflow
 - → Tracking object boundaries and verifying all memory accesses
- Temporal safety → e.g., use-after-free
 - → Tracking life time of objects and verifying all memory accesses

Idea in C/C++: if we implement everything *correctly*, we have opportunities to make the program much efficient (in terms of memory usage) and faster (in terms of execution speed)!

Error-prone C APIs: strncpy()

- strncpy(char *dest, const char *src, size_t n)

```
// BUG!  
char buf[BUFSIZ];  
strncpy(buf, input, sizeof(buf));
```

Error-prone C APIs: strncpy()

- strncpy(char *dest, const char *src, size_t n)

```
char buf[BUFSIZ];  
strncpy(buf, input, sizeof(buf) - 1);  
buf[sizeof(buf) - 1] = '\\0';
```


Error-prone C APIs: strncpy()

- NULL on the remaining bytes (why?)
- strncpy(buf, "A"*len(buf), buf) → leaving buf non-NULL-terminated
- dest and src should not be overlapping
- Return dest, not #chars copied!

```
char* strncpy(char *dest, const char *src, size_t n) {  
    size_t i;  
    for (i = 0; i < n && src[i] != '\0'; i++)  
        dest[i] = src[i];  
    for ( ; i < n; i++)  
        dest[i] = '\0';  
    return dest;  
}
```

Error-prone C APIs: strncat()

- dest is always NULL-terminated C-string
- Copy max $n + 1$ (w/ null)! \rightarrow strncat(dest, src, len - 1)
- Return dest, not #chars copied!

```
char* strncat(char *dest, const char *src, size_t n) {  
    size_t dest_len = strlen(dest);  
    for (size_t i = 0 ; i < n && src[i] != '\0' ; i++)  
        dest[dest_len + i] = src[i];  
    dest[dest_len + i] = '\0';  
    return dest;  
}
```

Suggestion for C-string manipulation

- Use: `snprintf(buf, sizeof(buf), ...)`
- Alternatives: `strncpy()`, `strlcat()`
 - Return #chars copied, or `strlen(dest)`
 - NULL-terminated, unless dest is full in cast of `strlcat()`

Outline

- Real-world examples:
 1. CVE-2017-15118: QEMU
 2. CVE-2014-4975: Wireshark
 3. CVE-2015-7547: glibc getaddrinfo)*
- Off-the-shelf defenses:
 1. Stack shield/canary (a.k.a, SSP in gcc)
 2. DEP (NX, W^X)
- Advance defenses: Shadow stack and CET

CVE-2017-15118: QEMU

```
1  #define NBD_MAX_NAME_SIZE 256
2
3  static int nbd_negotiate_handle_info(...) {
4      char name[NBD_MAX_NAME_SIZE + 1];
5      uint32_t namelen;
6
7      nbd_read(client->ioc, &namelen, sizeof(namelen), errp);
8      nbd_read(client->ioc, name, namelen, errp);
9  }
```

CVE-2014-4975: Ruby

```
1 void encodes(..) {
2     char buff[4096];
3     while (len >= 3) {
4         while (len >= 3 && sizeof(buff)-i >= 4) {
5             buff[i++] = ..; /* 4 times */;
6             s += 3;
7             len -= 3;
8         }
9         if (sizeof(buff)-i < 4) { /* flush */ }
10    }
11    if (len == 2) { ... }
12    else if (len == 1) { ... }
13
14    if (tail_lf) buff[i++] = '\n';
15    /* flush */
16 }
```

Then, How to Prevent Stack Overflow?

- Two approaches:
 - Bug prevention (i.e., correct bound checking)
 - Exploitation mitigation (i.e., making exploit harder)
1. Prevent the buffer overflow at the first place
 - (e.g., code analysis, designing better APIs)
 2. Protect “integrity” of ra, funcptr, etc (code pointers)
 - (e.g., exploitation mitigation → NX, canary)

Defense 1: Stack Canary

Stack Canary: Basic Idea

- Use a canary value as an indicator of the integrity of fp/ra
- Check the canary value right before using fp/ra (i.e., ret)

```
(top, growing)
|<--- current frame -->|<-- caller-->|
[buf                ][canary][fp][ra][...      ]
|<---- 16 ---->|
AAAAAAAAAAAAAAAAA XOXOXO BBBBCCCC... =>
                    (corrupted?)
```

Subtle Design Choices for the Stack Canary

1. Where to put? (e.g., right above ra? fp? local vars?)
2. Which value should I use? (e.g., secret? random? per exec? per func?)
3. How to check its integrity? (e.g., xor? cmp?)
4. What to do after you find corrupted? (e.g., crash? report?)

GCC's Stack Smashing Protector (SSP)

- Options: -fstack-protector/all/strong/explicit
 - Scope: all >> strong >> default >> explicit
 - e.g., use of alloca(), buffer, array, etc

```
$ cd lec01-stackovfl/ssp
$ cat ssp.c
$ make check-sspopts
...
```

Case Study: Using SSP in Linux (> 3.14)

- -fstack-protector:
 - 0.33% larger code size
 - 2.81% of the functions covered
- -fstack-protector-strong:
 - 2.4% larger code size
 - 20.5% of the functions covered

ref. <https://lwn.net/Articles/584225/>

SSP in Detail: Instrumentation

```
1  int func1_benign(int arg) { return arg; }
2
3  $ ./check-func.py ssp-explicit func1_benign
4  func1_benign()@ssp-explicit
5      push    rbp
6      mov     rbp, rsp
7      mov     DWORD PTR [rbp-0x4], edi
8      mov     eax, DWORD PTR [rbp-0x4]
9      pop     rbp
10     ret
```

SSP in Detail: Instrumentation

```
1  $ ./check-func.py ssp-all func1_benign
2  func1_benign()@ssp-all
3      push    rbp
4      mov     rbp, rsp
5      sub     rsp, 0x20
6      mov     DWORD PTR [rbp-0x14], edi
7      ! mov    rax, QWORD PTR fs:0x28      // read canary @TLS
8      ! mov    QWORD PTR [rbp-0x8], rax    // put it right above fp
9      ! xor    eax, eax                    // clear it off
10     mov     eax, DWORD PTR [rbp-0x14]
11     ! mov    rdx, QWORD PTR [rbp-0x8]    // fetch canary on stack
12     ! xor    rdx, QWORD PTR fs:0x28     // compare it with @TLS
13     ! je     func1_benign+0x31
14     ! call   __stack_chk_fail@plt       // stack smashed!
15     leave
16     ret
```

SSP in Detail: In Action

- Any interesting byte in canary?

```
$ ./canary
0x7ffd4d8e4278: 0xa                // saved arg
0x7ffd4d8e4280: 0x100000000        // dummy
0x7ffd4d8e4288: 0xdace23a197bb3d00 // canary
0x7ffd4d8e4290: 0x7ffd4d8e42c0     // fp
0x7ffd4d8e4298: 0x55d3082a51fc     // ra
```

About “Terminator” Canary

- Why is the terminator canary special?
 - 0x0d000aff: NULL (0x00), CR (0x0d), LF (0x0a) and EOF (0xff)
- SSP: Used to contain NULL/EOF/LF (06/2014, see [commit](#))
- SSP: Only contains NULL (@MSB) in a recent version

SSP: `__stack_chk_fail()`

- Immediately abort the program like below:

```
$ cd lec01-stackovfl/ssp
$ ./ovfl
*** stack smashing detected ***: ./ovfl terminated
Aborted
```

SSP: `__stack_chk_fail()` Implementation

```
1  // @debug/stack_chk_fail.c
2  void __attribute__((noreturn))
3  __stack_chk_fail (void) {
4      __fortify_fail ("stack smashing detected");
5  }
6
7  void __attribute__((noreturn))
8  __fortify_fail (const char *msg) {
9      __libc_message (2, "*** %s ***: %s terminated\n",
10                      msg, __libc_argv[0] ?: "<unknown>");
11 }
```

SSP: Security Issue in `__stack_chk_fail()`

- [CVE-2010-3192](#): Arbitrary read after stack smashing
 - `__libc_argv[0]` is under control
 - Breaking confidentiality, e.g., leaking private keys

SSP: New Implementation

```
1  /* Don't pass down __libc_argv[0] if we aren't doing bac  
2     since __libc_argv[0] may point to the corrupted stack  
3  __libc_message (need_backtrace ?  
4                  (do_abort | do_backtrace) : do_abort,  
5                  "*** %s ***: %s terminated\n",  
6                  msg,  
7                  (need_backtrace && __libc_argv[0] != NULL  
8                  ? __libc_argv[0] : "<unknown>"));
```

SSP: Placing Local Variables

```
long var1[32] = {1, };  
int (*var2)(const char *) = system;  
long var3 = 3;
```

```
$ cd lec01-stackovfl/ssp
```

```
$ make check-loc
```

```
func1_benign():
```

```
0x7ffd6375c580: 0x1 (var1)
```

```
0x7ffd6375c588: 0x2 (var2)
```

```
0x7ffd6375c590: 0x3 (var3)
```

```
...
```

```
func5_buf_and_funcptr():
```

```
0x7ffd6375c480: 0x7f5a62ecf380 (var2)
```

```
0x7ffd6375c488: 0x3 (var3)
```

```
0x7ffd6375c490: 0x1 (var1)
```

Limitation of Canary-based Approaches

1. Unprotected local variables (e.g., index, func ptrs)
2. Incrementally overwriting one byte at a time (in remote, fork())
3. Leaked canary (per execution)

Defense 2: DEP (NX, W^X)

- Data Execution Prevention (DEP)
 - aka, Non eXecutable, Writable ^ eXecutable
 - Basically, don't make writable region executable at the same time

```
cat /proc/self/maps
5606bdf09000-5606bdf0b000 r--p /usr/bin/cat
5606bdf0b000-5606bdf0f000 r-xp /usr/bin/cat
5606bdf13000-5606bdf14000 rw-p /usr/bin/cat
...
5606bef45000-5606bef66000 rw-p [heap]
7ffcd93c8000-7ffcd93ea000 rw-p [stack]
7ffcd93f7000-7ffcd93fa000 r--p [vvar]
7ffcd93fa000-7ffcd93fc000 r-xp [vdso]
```

Advance Defense 1: Shadow Stack

- Option: -fsanitize=shadow-call-stack
 - Replicate return addresses in a safe place, so called shadow stack
 - The shadow stack directly indicates modification of return addresses

```

      (top) +-----+
            |               v XXXX
<stack>   : [buf][canary][fp][ra][var][canary][fp][ra] ... (@rsp)
<shadow>  :               [ra]               [ra] ... (@gs)

```


Advance Defense 1: Shadow Stack

```
void vuln(char *arg) { char buf[32]; ... }
```

```
$ cd lec01-stackovfl/safestack
$ make check-shadowstack
-vuln()@shadowstack-no
+vuln()@shadowstack-yes
+  mov     r10,QWORD PTR [rsp]
+  xor     r11,r11
+  add     QWORD PTR gs:[r11],0x8
+  mov     r11,QWORD PTR gs:[r11]
+  mov     QWORD PTR gs:[r11],r10
+  push    rbp
+  mov     rbp,rsp
+  sub     rsp,0x40
+  ...
```

Advance Defense 2: Safe Stack

- Option: -fsanitize=safe-stack
- Two stacks: safe/unsafe stacks for sensitive/non-sensitive data
 - Preventing stack overflow to the sensitive data
 - Disentangling the leakage of stack pointers

```

                (top)
<stack>      : [buf][canary][fp][ra][var][canary][fp][ra] ... (@rsp)

<safe>       : [fp][ra][fp][ra] ... (@rsp)
<unsafe>     : [buf][canary][var][canary] ... (@fs)
               =====> overflow doesn't affect fp/ra + other sensitive data

```

Advance Defense 2: Safe Stack

```
void vuln(char *arg) { char buf[32]; ... }
```

```
$ cd lec01-stackovfl/safestack
```

```
$ make check-safestack
```

```
-vuln()@safestack-no
```

```
+vuln()@safestack-yes
```

```
    push    rbp
```

```
    mov     rbp, rsp
```

```
-   sub     rsp, 0x40
```

```
+   sub     rsp, 0x20
```

```
+   mov     rax, QWORD PTR [rip+0x8271] ; read the base of stacktop
```

```
+   mov     rcx, QWORD PTR fs:[rax]      ; fetch stacktop
```

```
+   mov     rdx, rcx
```

```
+   add     rdx, 0xfffffffffffffd0      ; allocate
```

```
+   mov     QWORD PTR fs:[rax], rdx     ; update stacktop
```

References

- [Smashing The Stack For Fun And Profit](#)
- [Scraps of notes on remote stack overflow exploitation](#)
- [Bypassing StackShield](#)
- [Bypassing Safe Stack](#)