

Lec01: Stack Overflow and Protections

Taesoo Kim

Goals and Lessons

- Learn about the **stack overflow** bugs
- Understand their security implications (i.e., control hijacking)
- Understand the off-the-shelf mitigation (i.e., ssp, DEP, RELO, ASLR)
- Learn them from the real-world examples (i.e., qemu/ruby/wireshark)

“Smashing The Stack For Fun And Profit”

CS101: What's Wrong?

```
1 | main() {  
2 |     char buf[16];  
3 |     scanf("%s", buf);  
4 | }
```

CS101: How to Fix?

```
1 | main() {  
2 |     char buf[16];  
3 |     scanf("%s", buf); // BUG!  
4 | }
```

1. `scanf("%s", &buf)`
2. `scanf("%16s", buf)`
3. `scanf("%15s", buf)`
4. `scanf("%as", &bufptr)`

Error-prone C APIs: scanf()

- Answer: scanf("%15s", buf)!

```
$ cd lec01-stackovfl/apis  
$ cat scanf.c  
$ make  
$ ./scanf  
...
```

Error-prone C APIs: scanf()

- `scanf("%16s", buf) // BUG!`

```
$ ./scanf
aaaaaaaaaaaaaaaaaaaaaaaaaaaaa
aaaaaaaaaaaaaaaaaaaaa
01: 61 (a)
02: 61 (a)
...
0e: 61 (a)
0f: 61 (a)
10: 61 (a)
11: 00 ()
12: BE ()
13: AD ()
14: DE ()
```

Security Implication: Control Hijacking

- The return address can be overwritten by an attacker's input
- Lead to control hijacking attacks (arbitrary execution)!

```
1 | main() {  
2 |     char buf[16];  
3 |     scanf("%s", buf); // BUG!  
4 | }
```


Basic Idea: Stack Smashing Attack

```

1  main() {
2      char buf[16];
3      scanf("%s", buf); // BUG!
4  }
```

(top, growing)

<stack>

(lower)

<==

libc_start_main()

|<-- caller-->|

[...][fp][ra] ...

| |

| +---> return address

+-----> frame pointer (**ebp**)

Basic Idea: Stack Smashing Attack

```

1 | main() {
2 |     char buf[16];
3 |     scanf("%s", buf); // BUG!
4 | }
```

(top, growing)		(lower)
	main()	libc_start_main()
<--- current frame -->	<-- caller-->	
[buf][fp][ra][...]] ...
<---- 16 --->		
	+----->	return address
	+----->	frame pointer (ebp)

Basic Idea: Stack Smashing Attack

```
1 | main() {  
2 |     char buf[16];  
3 |     scanf("%s", buf); // BUG!  
4 | }
```

(top, growing)

```
|<--- current frame -->|<-- caller-->|  
[buf                ][fp][ra][...      ]  
|<---- 16 ---->|  
AAAAAAAAAAAAAAAAABBBBCCCC... =>
```

!SEGFALT @eip=CCCC

Control Hijacking Attack: Where to Jump?

- Jump to the injected code (e.g., stack, environ, etc)

```
(top, growing)
|<--- current frame -->|<-- caller-->|
[buf                ][fp][ra][...      ]
|<----- 16 ---->|
AAAAAAAAAAAAAAAAABBBBXXXX[shellcode ..]
                    |      ^
                    +-----+ (1) attacker's input
```

Control Hijacking Attack: Where to Jump?

- Jump to the injected code (e.g., stack, environ, etc)

```
(top, growing)
|<--- current frame -->|
[buf                ][fp][ra] ... [SHELLCODE=...]
|<----- 16 ---->|           ^
AAAAAAAAAAAAAAAAABBBBXXXX   |
                        |     |
                        +-----+ (2) crafted environ
```

Control Hijacking Attack: Advanced Topics

1. Stack pivoting when frame pointer is crafted
2. Even off-by-one (e.g., writing NULL) is enough for stack pivoting

(top, growing)

|<--- current frame -->|<-- caller-->|

[buf [fp][ra][...]

|<---- 16 ---->|

1)AAAAAAAAAAAAAAAAAXXX -> enough to control (i.e., leave; *ret*)

2)AAAAAAAAAAAAAAAAAX -> off-by-one (e.g., scanf("%16s", buf))

Memory Safety in C/C++

- Spatial safety → e.g., buffer over/underflow
- Temporal safety → e.g., use-after-free

Addressing Memory Safety Issues in C/C++

- Spatial safety → e.g., buffer over/underflow
 - Tracking object boundaries and verifying all memory accesses
- Temporal safety → e.g., use-after-free
 - Tracking life time of objects and verifying all memory accesses

Idea in C/C++: if we implement everything *correctly*, we have opportunities to make the program much efficient (in terms of memory usage) and faster (in terms of execution speed)!

Error-prone C APIs: strcpy()/strncpy()



The strcpy() function copies the string pointed to by src, including the terminating null byte ('\0'), to the buffer pointed to by dest. The strncpy() function is similar, except that at most n bytes of src are copied.

```
char *strcpy(char *dest, const char *src);  
char *strncpy(char *dest, const char *src, size_t n);
```

```
1 | // BUG!  
2 | char buf[BUFSIZ];  
3 | strncpy(buf, input, sizeof(buf));
```

Error-prone C APIs: strncpy()

- Warning: dest might not be **null-terminated**!

```
1 | char buf[BUFSIZ];  
2 | strncpy(buf, input, sizeof(buf) - 1);  
3 | buf[sizeof(buf) - 1] = '\\0';
```

Error-prone C APIs: strncpy()

1. NULL on the remaining bytes (why?)!
2. strncpy(buf, "A"*len(buf), buf) → leaving buf non-NULL-terminated
3. Return dest, not #chars copied!

```
1 char* strncpy(char *dest, const char *src, size_t n) {  
2     size_t i;  
3     for (i = 0; i < n && src[i] != '\0'; i++)  
4         dest[i] = src[i];  
5     for ( ; i < n; i++) // Q1?  
6         dest[i] = '\0';  
7     return dest;  
8 }
```

Error-prone C APIs: strncat()

1. dest is always NULL-terminated C-string!
2. Copy max n + 1 (w/ null)! → strncat(dest, src, len - 1)
3. Return dest, not #chars copied!

```
1 char* strncat(char *dest, const char *src, size_t n) {  
2     size_t dest_len = strlen(dest);  
3     for (size_t i = 0 ; i < n && src[i] != '\0' ; i++)  
4         dest[dest_len + i] = src[i];  
5     dest[dest_len + i] = '\0';  
6     return dest;  
7 }
```

Suggestion for C-string Manipulation

- Use: `snprintf(buf, sizeof(buf), ...)`
- Alternatives: `strncpy()`, `strlcat()`
 - Return #chars copied, or `strlen(dest)`
 - The full size of dest (not the remaining space)
 - NULL-terminated, unless dest is full in cast of `strlcat()`

```
size_t strlcat(char *dst, const char *src, size_t size);  
size_t strncpy(char *dst, const char *src, size_t size);
```

Error-prone C APIs: fgets()

- Read at most *one less* than size!
- NULL-terminated!

```
char *fgets(char *s, int size, FILE *stream);  
$ cd lec01-stackovfl/apis  
$ ./fgets  
...
```

Error-prone C APIs: fgets()

- fgets() accepts input until it sees a newline (\n)
- It means that it accepts the terminator char: NULL!

```
char *fgets(char *s, int size, FILE *stream);
```

```
$ cd lec01-stackovfl/apis
```

```
$ echo -e "123\x0056" | ./fgets
```

```
01: 31 (1)
```

```
02: 32 (2)
```

```
03: 33 (3)
```

```
04: 00 ( )
```

```
05: 35 (5)
```

```
06: 36 (6)
```

```
07: 0A (\n)
```

Error-prone C APIs: fgets()

- size != strlen(input)

```
1 | size = fgets(input, sizeof(input), stdin);  
2 | // BUG!  
3 | if (strlen(input) < sizeof(dest)) {  
4 |     memcpy(dest, input, size);  
5 | }
```


Outline

- Real-world examples:
 1. Ex1. CVE-2017-15118: QEMU
 2. Ex2. CVE-2014-4975: Wireshark
 3. (Ex3. CVE-2015-7547: glibc getaddrinfo())*
- Off-the-shelf defenses:
 1. Stack shield/canary (a.k.a, SSP in gcc)
 2. DEP (NX, W^X)
- Advance defenses: Shadow/Safe Stack

CVE-2017-15118: QEMU NBD

- qemu-io f raw “nbd://localhost:10809/path” → looking up “path”

```
1  #define NBD_MAX_NAME_SIZE 256
2
3  static int nbd_negotiate_handle_info(...) {
4      char name[NBD_MAX_NAME_SIZE + 1];
5      uint32_t namelen;
6
7      nbd_read(client->ioc, &namelen, sizeof(namelen), errp);
8      nbd_read(client->ioc, name, namelen, errp);
9  }
```

CVE-2014-4975: Ruby

- ["a"*3070].pack("m4000") → encode(var, "aaa..", 3070, ..., true)

```
1 void encodes(VALUE str, const char *s, long len ...) {
2     char buff[4096];
3     while (len >= 3) {
4         while (len >= 3 && sizeof(buff)-i >= 4) {
5             buff[i++] = ...; /* 4 times */;
6             s += 3; len -= 3;
7         }
8         if (sizeof(buff)-i < 4) { /* flush */ }
9     }
10    if (len == 2) { buff[i++] = ...; /* 4 times */ }
11    else if (len == 1) { buff[i++] = ...; /* 4 times */ }
12    if (tail_lf) buff[i++] = '\n';
13    /* flush */
14 }
```

CVE-2014-4975: Ruby

- `["a"*3070].pack("m4000")` → `encode(var, "aaa..", 3070, ..., true)`
 - `buff` is used upto $3069 / 3 * 4 = 4092$ bytes (*)
 - Since one more byte left (`len == 1`), 4 more bytes are used (**)
 - `tail_lf` → one more byte: overflowing `"\n"`

```
1 void encodes(VALUE str, const char *s, long len ...) {
2     char buff[4096];
3     * while (len >= 3) {...}      // i -> 3069/3 x 4 = 4092 byte
4     if (len == 2) {...}
5     ** else if (len == 1) {...}  // i += 4 -> 4096 bytes
6     *** if (tail_lf)
7         buff[i++] = '\n'        // i += 1 (off-by-one)!
8 }
```

Then, How to Prevent Stack Overflow?

- Two approaches:
 - Bug prevention (i.e., correct bound checking)
 - Exploitation mitigation (i.e., making exploit harder)
1. Prevent the buffer overflow at the first place
 - (e.g., code analysis, designing better APIs)
 2. Protect “integrity” of ra, funcptr, etc (code pointers)
 - (e.g., exploitation mitigation → NX, canary)

Defense 1: Stack Canary

Stack Canary: Basic Idea

- Use a canary value as an indicator of the integrity of fp/ra
- Check the canary value right before using fp/ra (i.e., ret)

```
(top, growing)
|<--- current frame -->|<-- caller-->|
[buf                ][canary][fp][ra][...          ]
|<---- 16 ---->|
AAAAAAAAAAAAAAAAA XOXOXO BBBBCCCC... =>
                    (corrupted?)
```

Subtle Design Choices for the Stack Canary

1. Where to put? (e.g., right above ra? fp? local vars?)
2. Which value should I use? (e.g., secret? random? per exec? per func?)
3. How to check its integrity? (e.g., xor? cmp?)
4. What to do after you find corrupted? (e.g., crash? report?)

GCC's Stack Smashing Protector (SSP)

- Options: -fstack-protector/all/strong/explicit
 - Scope: all >> strong >> default >> explicit
 - e.g., use of `alloca()`, `buffer`, `array`, etc

```
$ cd lec01-stackovfl/ssp
$ cat ssp.c
$ make check-sspopts
...
```

Case Study: Using SSP in Linux (> 3.14)

- -fstack-protector:
 - 0.33% larger code size
 - 2.81% of the functions covered
- -fstack-protector-strong:
 - 2.4% larger code size
 - 20.5% of the functions covered

ref. <https://lwn.net/Articles/584225/>

SSP in Detail: Instrumentation

```
1  int func1_benign(int arg) { return arg; }
2
3  $ ./check-func.py ssp-explicit func1_benign
4  func1_benign()@ssp-explicit
5      push    rbp
6      mov     rbp, rsp
7      mov     DWORD PTR [rbp-0x4], edi
8      mov     eax, DWORD PTR [rbp-0x4]
9      pop     rbp
10     ret
```

SSP in Detail: Instrumentation

```
1  $ ./check-func.py ssp-all func1_benign
2  func1_benign()@ssp-all
3      push    rbp
4      mov     rbp, rsp
5      sub     rsp, 0x20
6      mov     DWORD PTR [rbp-0x14], edi
7      ! mov    rax, QWORD PTR fs:0x28      // read canary @TLS
8      ! mov    QWORD PTR [rbp-0x8], rax    // put it right above fp
9      ! xor    eax, eax                    // clear it off
10     mov     eax, DWORD PTR [rbp-0x14]
11     ! mov    rdx, QWORD PTR [rbp-0x8]    // fetch canary on stack
12     ! xor    rdx, QWORD PTR fs:0x28     // compare it with @TLS
13     ! je     func1_benign+0x31
14     ! call   __stack_chk_fail@plt      // stack smashed!
15     leave
16     ret
```

SSP in Detail: In Action

- Any interesting byte in canary?

```
$ ./canary
0x7ffd4d8e4278: 0xa                // saved arg
0x7ffd4d8e4280: 0x100000000        // dummy
0x7ffd4d8e4288: 0xdace23a197bb3d00 // canary
0x7ffd4d8e4290: 0x7ffd4d8e42c0     // fp
0x7ffd4d8e4298: 0x55d3082a51fc     // ra
```

About “Terminator” Canary

- Why is the terminator canary special?
 - 0x0d000aff: NULL (0x00), CR (0x0d), LF (0x0a) and EOF (0xff)
- SSP: Used to contain NULL/EOF/LF (06/2014, see [commit](#))
- SSP: Only contains NULL (@LSB) in a recent version

SSP: __stack_chk_fail()

- Immediately abort the program like below:

```
! mov    rdx,QWORD PTR [rbp-0x8]    // fetch canary on stack
! xor    rdx,QWORD PTR fs:0x28      // compare it with @TLS
! je     func1_benign+0x31
! call   __stack_chk_fail@plt         // stack smashed!
```

```
$ cd lec01-stackovfl/ssp
```

```
$ ./ovfl
```

```
*** stack smashing detected ***: ./ovfl terminated
```

```
Aborted
```

SSP: `__stack_chk_fail()` Implementation

```
1  // @debug/stack_chk_fail.c
2  void __attribute__((noreturn))
3  __stack_chk_fail (void) {
4      __fortify_fail ("stack smashing detected");
5  }
6
7  void __attribute__((noreturn))
8  __fortify_fail (const char *msg) {
9      __libc_message (2, "*** %s ***: %s terminated\n",
10                      msg, __libc_argv[0] ?: "<unknown>");
11 }
```


SSP: Security Issue in `__stack_chk_fail()`

- [CVE-2010-3192](#): Arbitrary read after stack smashing
 - `__libc_argv[0]` is under control
 - Breaking confidentiality, e.g., leaking private keys

SSP: New Implementation

```
1  /* Don't pass down __libc_argv[0] if we aren't doing  
2     backtrace since __libc_argv[0] may point to the  
3     corrupted stack. */  
4  __libc_message (need_backtrace ?  
5                  (do_abort | do_backtrace) : do_abort,  
6                  "*** %s ***: %s terminated\n",  
7                  msg,  
8                  (need_backtrace && __libc_argv[0] != NULL  
9                  ? __libc_argv[0] : "<unknown>"));
```

SSP: Placing Local Variables

```
long var1[32] = {1, };  
int (*var2)(const char *) = system;  
long var3 = 3;
```

```
$ cd lec01-stackovfl/ssp
```

```
$ make check-loc
```

```
func1_benign():
```

```
    0x7ffd6375c580: 0x1 (var1)
```

```
    0x7ffd6375c588: 0x2 (var2)
```

```
    0x7ffd6375c590: 0x3 (var3)
```

```
...
```

```
func5_buf_and_funcptr():
```

```
    0x7ffd6375c480: 0x7f5a62ecf380 (var2)
```

```
    0x7ffd6375c488: 0x3 (var3)
```

```
    0x7ffd6375c490: 0x1 (var1)
```

Limitation of Canary-based Approaches

1. Unprotected local variables (e.g., index, func ptrs)
2. Incrementally overwriting one byte at a time (in remote, fork())
3. Leaked canary (per execution)

Defense 2: DEP (NX, W^X)

- Data Execution Prevention (DEP)
 - aka, Non eXecutable, Writable ^ eXecutable
 - Basically, don't make writable region executable at the same time

```
$ cat /proc/self/maps
5606bdf09000-5606bdf0b000 r--p /usr/bin/cat
5606bdf0b000-5606bdf0f000 r-xp /usr/bin/cat
5606bdf13000-5606bdf14000 rw-p /usr/bin/cat
...
5606bef45000-5606bef66000 rw-p [heap]
7ffcd93c8000-7ffcd93ea000 rw-p [stack]
7ffcd93f7000-7ffcd93fa000 r--p [vvar]
7ffcd93fa000-7ffcd93fc000 r-xp [vdso]
```

Advance Defense 1: Shadow Stack

- Option: -fsanitize=shadow-call-stack
 - Replicate return addresses in a safe place, so called shadow stack
 - The shadow stack directly indicates modification of return addresses

```

      (top) +-----+
            |               v XXXX
<stack>   : [buf][canary][fp][ra][var][canary][fp][ra] ... (@rsp)
<shadow>  :               [ra]               [ra] ... (@gs)

```

Advance Defense 1: Shadow Stack

```
void vuln(char *arg) { char buf[32]; ... }
```

```
$ cd lec01-stackovfl/safestack
$ make check-shadowstack
-vuln()@shadowstack-no
+vuln()@shadowstack-yes
+  mov     r10,QWORD PTR [rsp]
+  xor     r11,r11
+  add     QWORD PTR gs:[r11],0x8
+  mov     r11,QWORD PTR gs:[r11]
+  mov     QWORD PTR gs:[r11],r10
+  push    rbp
+  mov     rbp,rsp
+  sub     rsp,0x40
+  ...
```

Advance Defense 2: Safe Stack

- Option: -fsanitize=safe-stack
- Two stacks: safe/unsafe stacks for sensitive/non-sensitive data
 - Preventing stack overflow to the sensitive data
 - Disentangling the leakage of stack pointers

```

                (top)
<stack>      : [buf][canary][fp][ra][var][canary][fp][ra] ... (@rsp)

<safe>       : [fp][ra][fp][ra] ... (@rsp)
<unsafe>     : [buf][canary][var][canary] ... (@fs)
               =====> overflow doesn't affect fp/ra + other sensitive data

```


Advance Defense 2: Safe Stack

```
void vuln(char *arg) { char buf[32]; ... }
```

```
$ cd lec01-stackovfl/safestack
```

```
$ make check-safestack
```

```
-vuln()@safestack-no
```

```
+vuln()@safestack-yes
```

```
    push    rbp
```

```
    mov     rbp, rsp
```

```
-   sub     rsp, 0x40
```

```
+   sub     rsp, 0x20
```

```
+   mov     rax, QWORD PTR [rip+0x8271] ; read the base of stacktop
```

```
+   mov     rcx, QWORD PTR fs:[rax]      ; fetch stacktop
```

```
+   mov     rdx, rcx
```

```
+   add     rdx, 0xfffffffffffffd0      ; allocate
```

```
+   mov     QWORD PTR fs:[rax], rdx     ; update stacktop
```

Summary

- **Stack overflow** vulnerabilities
- Error-prone APIs: `strncpy`, `scanf`, `fgets`, etc.
- Understand its security implication (via stack smashing)
- Mitigation schemes: stack canary and DEP

References

- [Smashing The Stack For Fun And Profit](#)
- [Scraps of notes on remote stack overflow exploitation](#)
- [Bypassing StackShield](#)
- [Bypassing Safe Stack](#)