


Lec04: Heap-related Vulnerabilities

Taesoo Kim

Goals and Lessons

- Learn about the **heap**-related vulnerabilities
 - Buffer overflow/underflow, out-of-bound read
 - **Use-after-free**, including double frees
- Understand their security implications
- Learn them from the real-world examples

Trends of Vulnerability Classes



Ref. [Exploitation Trends: From Potential Risk to Actual Risk, RSA 2015](#)

Classifying Heap Vulnerabilities

- Common: buffer overflow/underflow, out-of-bound read
 - *Much prevalent* (i.e., quality, complexity)
 - *Much critical* (i.e., larger attack surface)
- Heap-specific issues:
 - **Use-after-free** (e.g., dangled pointers)
 - Incorrect uses (e.g., double frees)

Simple High-level Interfaces

```
// allocate a memory region (an object)  
void *malloc(size_t size);  
// free a memory region  
void free(void *ptr);  
  
// allocate a memory region for an array  
void *calloc(size_t nmemb, size_t size);  
// resize/reallocate a memory region  
void *realloc(void *ptr, size_t size);  
  
// new Type == malloc(sizeof(Type))  
// new Type[size] == malloc(sizeof(Type)*size)
```

CS101: Heap Allocators

Q0. `ptr = malloc(size); *ptr?`

Q1. `ptr = malloc(0); ptr == NULL?`

Q2. `ptr = malloc(-1); ptr == NULL?`

Q3. `ptr = malloc(size); ptr == NULL but valid? /* vaddr = 0 */`

Q4. `free(ptr); ptr == NULL?`

Q5. `free(ptr); *ptr?`

Q6. `free(NULL)?`

Q7. `realloc(ptr, size); ptr valid?`

Q8. `ptr = calloc(nmemb, size); *ptr?`

CS101: Common Goals of Heap Allocators

1. Performance
2. Memory fragmentation
3. (sometime) Security

// either fast, secure, (external) fragmentation!

1. malloc() -> mmap()	& free() -> unmap()
2. malloc() -> brk()	& free() -> nop
3. malloc() -> base += size; return base	& free() -> nop

Memory Allocators

Allocators	B	I	C	Description (applications)
ptmalloc	✓	✓	✓	A default allocator in Linux
dlmalloc	✓	✓	✓	An allocator that ptmalloc is based on
jemalloc	✓		✓	A default allocator in FreeBSD
tcmalloc	✓	✓	✓	A high-performance allocator from Google
PartitionAlloc	✓		✓	A default allocator in Chromium
libumem	✓		✓	A default allocator in Solaris

Common Design Choices (Security-Related)


1. **Binning:** size-base groups/operations
 - e.g., caching the same size objects together
2. **In-place metadata:** metadata before/after or even inside
 - e.g., putting metadata inside the freed region
3. **Cardinal metadata:** no encoding, direct pointers and sizes
 - e.g., using raw pointers for linked lists

Memory Allocators

Allocators	B	I	C	Description (applications)
ptmalloc	✓	✓	✓	A default allocator in Linux
dlmalloc	✓	✓	✓	An allocator that ptmalloc is based on
jemalloc	✓		✓	A default allocator in FreeBSD
tcmalloc	✓	✓	✓	A high-performance allocator from Google
PartitionAlloc	✓		✓	A default allocator in Chromium
libumem	✓		✓	A default allocator in Solaris

ptmalloc in Linux: Memory Allocation


```
ptr = malloc(size);
```



ptmalloc in Linux: Data Structure

```
struct malloc_chunk {  
    // size of "previous" chunk  
    // (only valid when the previous chunk is freed, P=0)  
    size_t prev_size;  
  
    // size in bytes (aligned by double words): lower bits  
    // indicate various states of the current/previous chunk  
    //   A: allocated in a non-main arena  
    //   M: mmaped  
    //   P: "previous" in use (i.e., P=0 means freed)  
    size_t size;  
  
    [...]  
};
```

ptmalloc in Linux: Memory Allocation



(a) allocated chunk

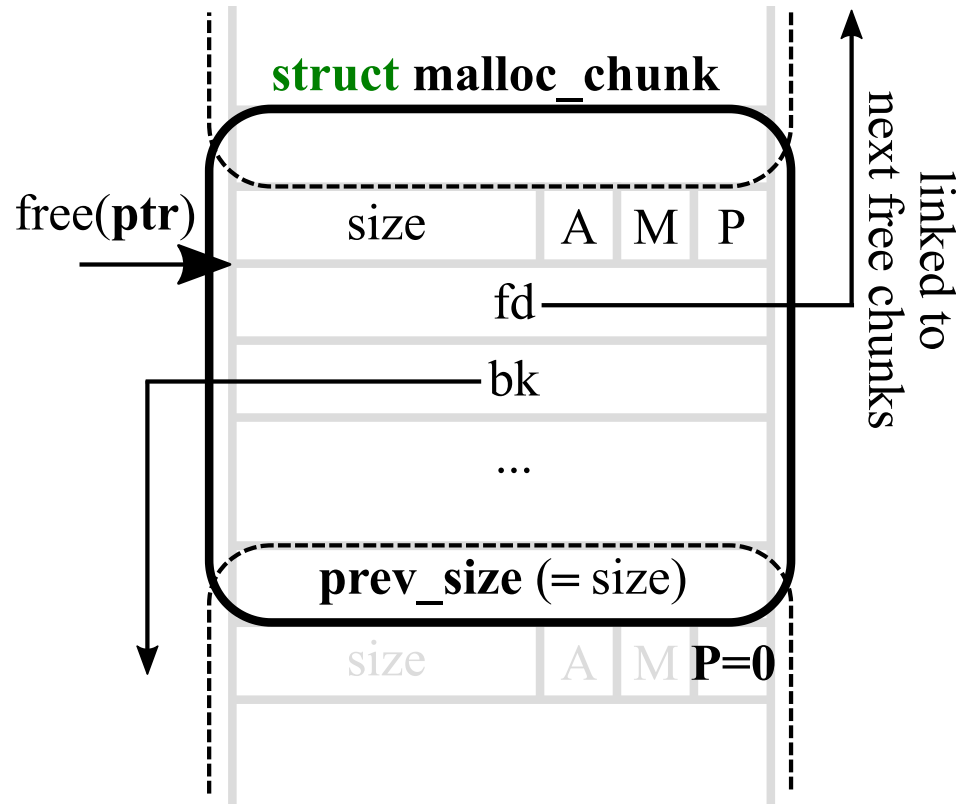
Remarks: Memory Allocation

- Given a allocated ptr,
 1. Immediately lookup its size!
 2. Check if the previous object is allocated/freed ($P = 0$ or 1)
 3. Iterate to the next object (not previous object if allocated)
 4. Check if the next object is allocated/freed (the next, next one's P)

ptmalloc in Linux: Data Structure

```
struct malloc_chunk {  
    [...]  
    // double links for free chunks in small/large bins  
    // (only valid when this chunk is freed)  
    struct malloc_chunk* fd;  
    struct malloc_chunk* bk;  
  
    // double links for next larger/smaller size in largebins  
    // (only valid when this chunk is freed)  
    struct malloc_chunk* fd_nextsize;  
    struct malloc_chunk* bk_nextsize;  
};
```

ptmalloc in Linux: Memory Free




(b) free chunk
(e.g., small bin)

Remarks: Memory Free

- Given a free-ed ptr,
 1. All benefits as an allocated ptr (previous remarks)
 2. Iterate to previous/next free objects via fd/bk links
- Invariant: **no two adjacent** free objects ($P = 0$)
 1. When free(), check if previous/next objects are free and consolidate!

Understanding Modern Heap Allocators

- Maximize memory usage: using free memory regions!
- Data structure to minimize fragmentation (i.e., fd/bk consolidation)
- Data structure to maximize performance (i.e., $O(1)$ in free/malloc)



Security Implication of Heap Overflows

- All metadata can be easily modified/crafted!
- Or even new alloc/free objects are created for benefits (and fun!)

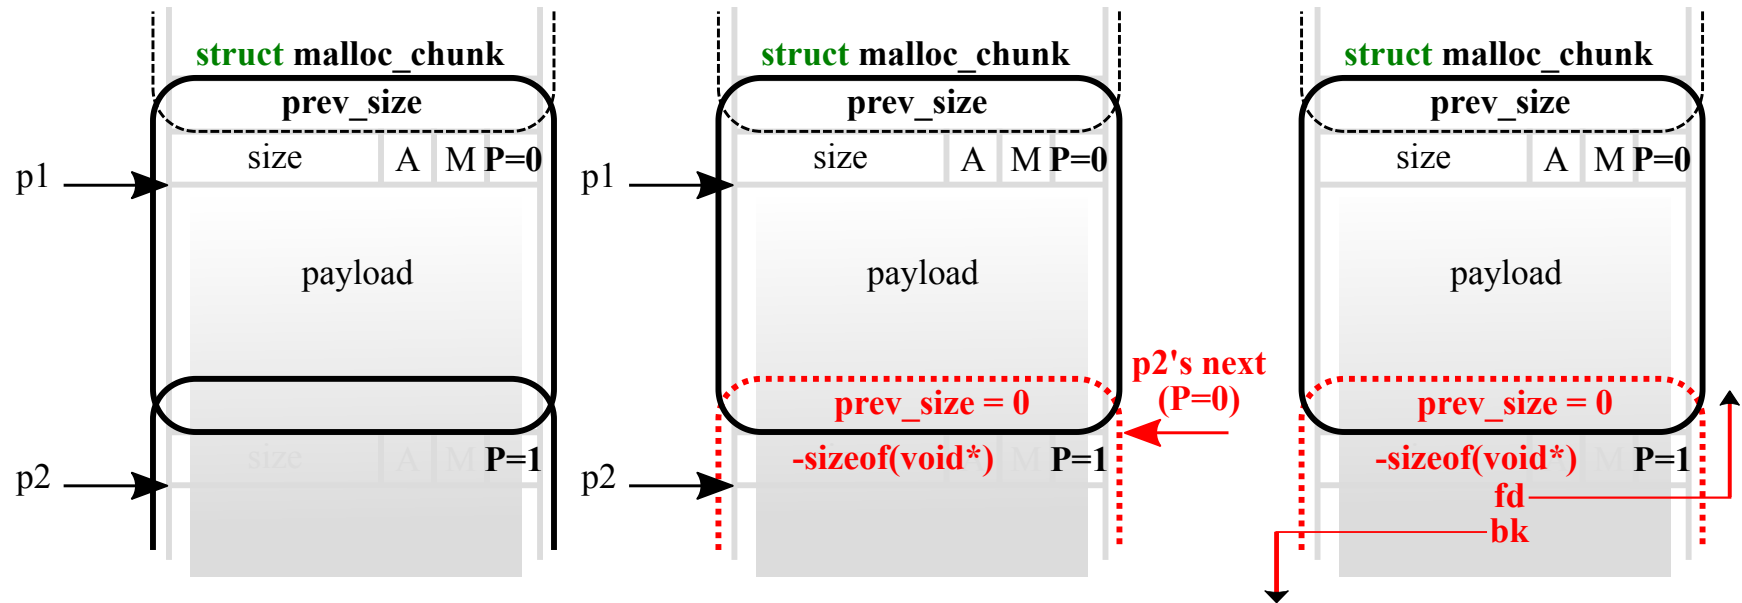
```
void *p1 = malloc(sz);  
void *p2 = malloc(sz);
```

```
/* overflow on p1 */
```

```
free(p1);
```

Example: Unsafe Unlink (< glibc 2.3.3)

1. Overwriting to p2's size to `-sizeof(void*)`, treating now as if p2 is free
2. When `free(p1)`, attempt to consolidate it with p2 as p2 is free



Example: Unsafe Unlink (< glibc 2.3.3)

- To consolidate, perform unlike on p2 (removing p2 from the linked list)
- Crafted fd/bk when unlink() result in an arbitrary write!

```
p2's fd = dst - offsetof(struct malloc_chunk, bk);  
p2's bk = val;
```

```
-> *dst = val (arbitrary write!)
```

```
#define unlink(P, BK, FD)
```

```
    FD = P->fd;
```

```
    BK = P->bk;
```

```
    FD->bk = BK;
```

```
    BK->fd = FD;
```

```
    ...
```

Example: Mitigation on Unlink (glibc 2.27)

```
#define unlink(AV, P, BK, FD)
    /* (1) checking if size == the next chunk's prev_size */
    * if (chunksize(P) != prev_size(next_chunk(P)))
    *     malloc_printerr("corrupted size vs. prev_size");
    *     FD = P->fd;
    *     BK = P->bk;
    *     /* (2) checking if prev/next chunks correctly point to me */
    *     if (FD->bk != P || BK->fd != P)
    *         malloc_printerr("corrupted double-linked list");
    *     else {
    *         FD->bk = BK;
    *         BK->fd = FD;
    *         ...
    *     }
```

Heap Exploitation Techniques!

Fast bin dup	House of einherjar
Fast bin dup into stack	House of orange
Fast bin dup consolidate	Tcache dup
Unsafe unlink	Tcache house of spirit
House of spirit	Tcache poisoning
Poison null byte	Tcache overlapping chunks
House of lore	*Unsorted bin into stack
Overlapping chunks 1	*Fast bin into other bin
Overlapping chunks 2	*Overlapping small chunks
House of force	*Unaligned double free
Unsorted bin attack	*House of unsorted einherjar

NOTE. * are what our group recently found and reported!

Use-after-free

- Simple in concept, but difficult to spot in practice!
- Why is it so critical in terms of security?

```
1 | int *ptr = malloc(size);  
2 | free(ptr);  
3 |  
4 | *ptr; // BUG. use-after-free!
```


Use-after-free

1. What would be the `*ptr`? if nothing happened?
2. What if another part of code invoked `malloc(size)`?


```
1 | int *ptr = malloc(size);  
2 | free(ptr);  
3 |  
4 | *ptr; // BUG. use-after-free!
```

Use-after-free: Security Implication

1. What would be the *ptr? if nothing happened?
 - → Heap pointer leakage (e.g., fd/bk)
2. What if another part of code invoked malloc(size)?
 - → Hijacking function pointers (e.g., handler)


```
1  struct msg { void (*handler)(); };  
2  
3  struct msg *ptr = malloc(size);  
4  free(ptr);  
5  // ...?  
6  ptr->handler(); // BUG. use-after-free!
```

Use-after-free with Application Context




```
1 | class Div: Element;  
2 | class Body: Element;  
3 | class Document { Element* child; };
```

Use-after-free with Application Context



```
1 class Div: Element;
2 class Body: Element;
3 class Document { Element* child; };
4
5 // (a) memory allocations
6 Document *doc = new Document();
7 Body *body = new Body();
8 Div *div = new Div();
```

Dangled Pointers and Use-after-free



```
1 // (b) using memory: propagating pointers
2 doc->child = body;
3 body->child = div;
4
5 // (c) memory free: doc->child is now dangled
6 delete body;
7
8 // (d) use-after-free: dereference the dangled pointer
9 if (doc->child)
10     doc->child->getAlign();
```

Double Free

1. What happen when free two times?
2. What happen for following malloc()s?

```
1 | char *ptr = malloc(size);  
2 | free(ptr);  
3 | free(ptr); // BUG!
```

Binning and Security Implication

- e.g., size-based caching (e.g., fastbin)

```
(fastbin)
  Bins
sz=16 [ -]--->[fd]--->[fd]-->NULL
sz=24 [ -]--->[fd]--->NULL
sz=32 [ -]--->NULL
...
```

Double Free

- Bins after doing free() two times

```

1 | char *ptr = malloc(sz=16);
2 | free(ptr);
3 | free(ptr); // BUG!

```

```

(fastbin)
      Bins  ptr      ptr
sz=16 [ -]--->[XX]--->[XX]--->[fd]--->[fd]-->NULL
sz=24 [ -]--->[fd]--->NULL
sz=32 [ -]--->NULL
      ...

```


Double Free: Security Implication

```

1 | char *ptr = malloc(sz=16);
2 | free(ptr);
3 | free(ptr); // BUG!
4 |
5 | ptr1 = malloc(sz=16) // hijacked!
6 | ptr2 = malloc(sz=16) // hijacked!

```

(fastbin)

Bins

```

          +-----+
          |         |
sz=16 [ -]--+ [XX]--->[XX] +-->[fd]--->[fd]-->NULL
sz=24 [ -]--->[fd]--->NULL
sz=32 [ -]--->NULL
...

```

Double Free: Mitigation

- Check if the bin contains the pointer that we'd like to free()

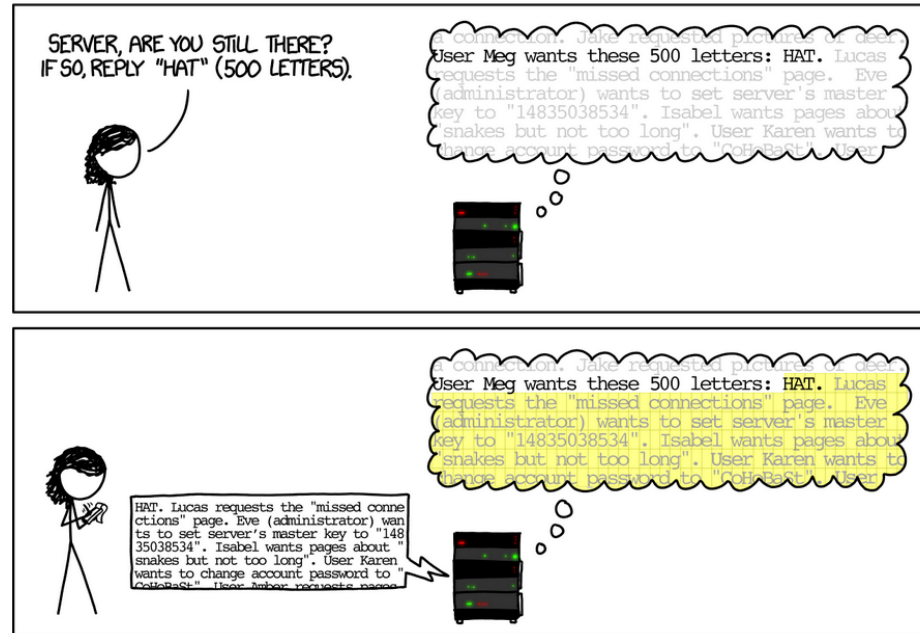
```
1  // @glibc/malloc/malloc.c
2
3  /* Check that the top of the bin is not the record we are
4     adding (i.e., double free). */
5  if (__builtin_expect (old == p, 0))
6     malloc_printerr ("double free or corruption (fasttop)"
7     ...
```

Exercise: Real-world Examples

- Ex1. OpenSSL (CVE-2014-0160)
- Ex2. Wireshark (CVE-2018-11360)
- Ex3. Linux vmcache (CVE-2018-17182)*

CVE-2014-0160: OpenSSL, Heartbleed

- Information leakage (i.e., private keys)



CVE-2014-0160: OpenSSL, Heartbleed

- “Heartbeat” messages to ensure the connection is alive

```
                |<--- len'' --->|  
-> req: [REQ][len'][payload .... ]  
<- res: [RES][len'][payload .... ][padding]
```

```
len' == len''?  
what if len' < len''?  
what if len' > len''?
```

CVE-2014-0160: OpenSSL, Heartbleed

```
1  /* Read type and payload length first */
2  hbtype = *p++;
3  n2s(p, payload);
4  pl = p;
5
6  if (hbtype == TLS1_HB_REQUEST) {
7      bp = OPENSSL_malloc(1 + 2 + payload + padding);
8
9      /* Enter response type, length and copy payload */
10     *bp++ = TLS1_HB_RESPONSE;
11     s2n(payload, bp);
12     memcpy(bp, pl, payload);
```

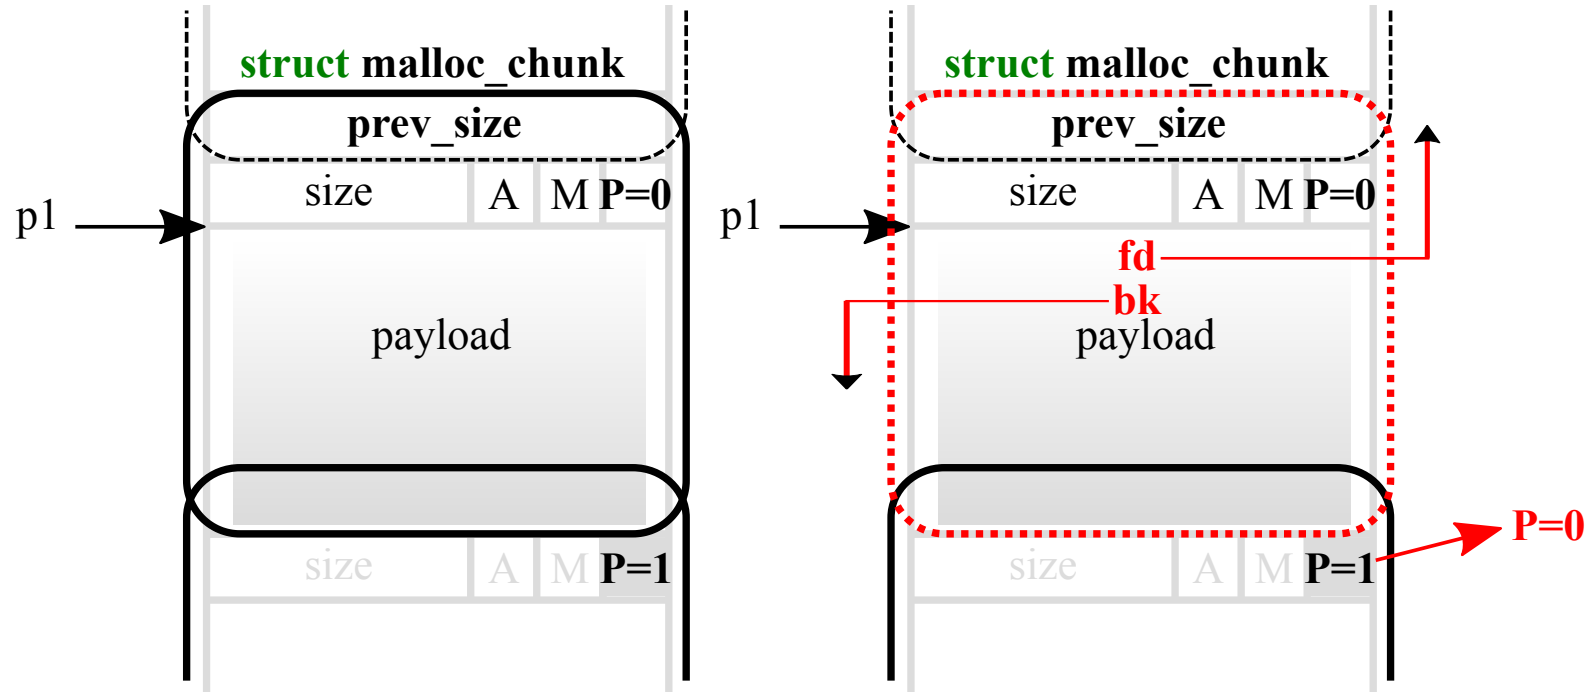
CVE-2014-0160: OpenSSL, Heartbleed

```
1      unsigned int payload;  
2      ...  
3      + /* Read type and payload length first */  
4      + if (1 + 2 + 16 > s->s3->rrec.length)  
5      +     return 0; /* silently discard */  
6      +  
7      + hbtype = *p++;  
8      + n2s(p, payload);  
9      + // NOTE. int overflow?  
10     + if (1 + 2 + payload + 16 > s->s3->rrec.length)  
11     +     return 0; /* silently discard per RFC 6520 sec. 4 */  
12     + pl = p;  
13     +  
14     + if (hbtype == TLS1_HB_REQUEST) { ... }
```

CVE-2018-11360: Wireshark

```
1  // NOTE. What's the semantics of data/len?
2  void IA5_7BIT_decode(unsigned char *dest,
3                        const unsigned char *src, int len) {
4      int i, j;
5      gunichar buf;
6
7      for (i = 0, j = 0; j < len; j++) {
8          buf = char_def_ia5_alphabet_decode(src[j]);
9          i += g_unichar_to_utf8(buf, &(dest[i]));
10     }
11     dest[i]=0;
12     return;
13 }
```


Security Implication of off-byte-one (NULL)



CVE-2018-17182: Linux vmcache*

- An optimization path for the single thread
- `mm→vmcache_seqnum` wraps around by another thread
 - → Dangled pointers suddenly become valid!

```
1 void vmacache_flush_all(struct mm_struct *mm) {  
2     /* Single threaded tasks need not iterate the entire list  
3     * process. We can avoid the flushing as well since the mm  
4     * was increased and don't have to worry about other threa  
5     * seqnum. Current's flush will occur upon the next lookup  
6     if (atomic_read(&mm->mm_users) == 1)  
7         return;  
8     ...  
9 }
```

Summary

- Two classes of **heap**-related vulnerabilities
 - Traditional: buffer overflow/underflow, out-of-bound read
 - Specific: **use-after-free**, **dangled pointers**, double free
- Understand why they are security critical and non-trivial to eliminate!
- Mitigation approaches taken by allocators

References

- [CVE-2014-0160](#)
- [CVE-2018-11360](#)
- [CVE-2018-17182](#)
- [Vudo - An object superstitiously believed to embody magical powers](#)