

# Lec05: Advanced Topics in Security

*Taesoo Kim*

# Goals and Lessons

- Understand *modern* attack/defenses: **ROP** and **CFI**
- Understand three classes of vulnerabilities
  - **Type confusion**: e.g., bad casting in C++
  - **Race condition**: e.g., double fetching
  - **Uninitialized read**: e.g., struct padding issues
- Learn them from real-world examples (e.g., Chrome, Linux)

# Modern Exploit against DEP (NX)

- Return-oriented Programming (ROP)
- Reusing code snippets (called gadgets) instead of injecting shellcode
  - e.g., ret-to-libc: `ret` → `system("/bin/sh")`

# Example: Stack Smashing (w/o DEP)

```
1 | main() {  
2 |     char buf[16];  
3 |     scanf("%s", buf);  
4 | }
```

(top)

```
    [buf  ]  
||[ ... ]  
||[ra    ]---+  
||[ ... ]   |  
||[shell]<--+ (inject shellcode)  
||[code  ]  
vv[ ... ]  
    [    ]  
(w/o DEP)
```

# Example: Stack Smashing vis Ret-to-libc

```

1 | main() {
2 |     char buf[16];
3 |     scanf("%s", buf);
4 | }

```

(top)

[buf ]	[buf ]	
[ ... ]	[ ... ]	
[ra ]---	[ra ]---	system() in libc
[ ... ]	[dummy]	
[shell]<---	[arg1 ]---	"/bin/sh"
[code ]	[ ... ]	
vv[ ... ]	[ ... ]	
[ ]	[ ]	
(w/o DEP)	(ret-to-libc)	

# Example: Stack Smashing vis Ret-to-libc

- Q. What happens when `system()` returns?
- Q. Is there any way to “gracefully” terminate?

```
(top)
[buf  ]
[ ... ]
[ra   ]---> system() in libc
[dummy]
[arg1 ]---> "/bin/sh"
[ ... ]
[ ... ]
[     ]
(ret-to-libc)
```

# Example: Stack Smashing vis Ret-to-libc

- A. “dummy” → exit()
- A. its first argument → 0

```
(top)                                system("/bin/sh")
[buf  ]                               exit(0)
[ ... ]
[ra1  ]---> system() in libc
[ra2  ]---> exit()
[arg1 ]---> "/bin/sh"
[arg2 ]---> 0
[ ... ]
[      ]
(ret-to-libc)
```

# Example: Executing Two Functions

- The return address of main() is smashed by the buffer overflow
- It returns to system() instead of the original caller

```
main(): ret
```

```
(top)
    [buf  ]
    [ ... ]
esp =>[ra1  ]---> system() in libc
    [ra2  ]---> exit()
    [arg1  ]---> "/bin/sh"
    [arg2  ]---> 0
    [ ... ]
    [      ]
(ret-to-libc)
```



# Example: Executing Two Functions

- dummy (i.e., ptr to exit()) is now considered as caller of system()
- arg1 is the first argument of system, “/bin/sh”

```

main(): ret
system(): ..
  caller: exit()
  arg: "/bin/sh"

(top)
  [buf  ]
  [ ... ]
  [      ]
esp => [ra2 ]---> exit()
      [arg1 ]---> "/bin/sh"
      [arg2 ]---> 0
      [ ... ]
      [      ]
(ret-to-libc)

```

# Example: Executing Two Functions

- dummy (i.e., ptr to `exit()`) is now considered as caller of `system()`
- `arg1` is the first argument of `system`, `"/bin/sh"`

```
main(): ret
system(): ret
```

```
(top)
    [buf  ]
    [ ... ]
    [      ]
esp =>[ra2  ]---> exit()
    [arg1  ]---> "/bin/sh"
    [arg2  ]---> 0
    [ ... ]
    [      ]
(ret-to-libc)
```

# Example: Executing Two Functions

- dummy (i.e., ptr to exit()) is now considered as caller of system()
- arg1 is the first argument of system, “/bin/sh”

main(): ret	(top)
system(): ret	[buf ]
exit():	[ ... ]
caller: "/bin/sh" (!!)	[ ]
arg: 0	[ra2 ]---> exit()
	esp =>[arg1 ]---> "/bin/sh"
	[arg2 ]---> 0
	[ ... ]
	[ ]
	(ret-to-libc)

# Example: Execution More than Two Funcs?

- Can we chain three functions in this way? No!
- ROP generalizes this approach by using

```
(top)                                system("/bin/sh")
[buf  ]                               exit(0)
[ ... ]
[ra1  ]---> system() in libc
[ra2  ]---> exit()
[arg1 ]---> "/bin/sh"                <- "/bin/sh" vs func3?
[arg2 ]---> 0
[ ... ]
[     ]
(ret-to-libc)
```

# Example: ROP

- Cleaning up stacks by using pop/ret gadgets
- Chaining them in a general fashion!

```
main(): ret
system(): ret

(top)
[buf ]
[ ... ]
[     ]
** esp => [ra2 ]---> exit()
[arg1 ]---> "/bin/sh"
[arg2 ]---> 0
[ ... ]
[     ]
(ret-to-libc)
```

# Example: ROP

- Cleaning up stacks by using pop/ret gadgets
- Chaining them in a general fashion!

```

main(): ret
system(): ret

                (top)
                [buf  ]
                [ ... ]
                [      ]
**  esp => [ra2  ]----> pop/ret
                [arg1 ]----> "/bin/sh"
                [ ... ]
                [ ... ]
                [      ]
                (ROP)

```

# Example: ROP

- Cleaning up stacks by using pop/ret gadgets
- Chaining them in a general fashion!

```
main(): ret
system(): ret
pop
```

```
(top)
[buf  ]
[ ... ]
[      ]
[      ]
esp =>[arg1 ]---> "/bin/sh"
[ ra  ]---> exit()
[ ... ]
[arg2 ]----> 0
(ROP)
```

# Example: ROP

- Cleaning up stacks by using pop/ret gadgets
- Chaining them in a general fashion!

```

main(): ret
system(): ret
pop
ret
                                     (top)
                                     [buf  ]
                                     [ ... ]
                                     [     ]
                                     [     ]
                                     [     ]
                                     [     ]
esp => [ ra  ]----> exit()
                                     [ ... ]
                                     [arg2 ]----> 0
                                     (ROP)

```



# Example: Beyond Two Functions!

- Cleaning up stacks by using pop/ret gadgets
- Chaining them in a general fashion!

```

(top)                                system("/bin/sh")
[buf  ]                               exit(0)
[ ... ]                              func3(arg3, arg4)
[ra1  ]---> system() in libc |
* [gadgt]---> pop/ret         | system()
[arg1  ]---> "/bin/sh"       |
[ra2  ]---> exit()           =
* [gadgt]---> pop/ret         = exit()
[arg2  ]---> 0               =
[ra3  ]---> func3()          +
* [gadgt]---> pop/pop/ret    + func3()
[arg3  ] ...                 +

```

# Defenses: Control-flow Integrity


- Control-flow at the compilation time should be enforced at runtime!
- Vendors adoptions:
  - [Control-flow Guard \(CFG\) by Microsoft from Windows 10/8.1](#)
  - [LLVM forward CFI by Android \(Pixel 3\)](#)
  - [LLVM forward CFI by Google Chrome \(dev\)](#)
- Hardware solutions:
  - Intel: [Control-flow Enforcement \(CET\)](#)
  - ARM v8,3: [PAC in iOS 12 by Apple](#)

# Basic Idea: Enforcing Control-flow Graphs

- Forward CFI: protecting indirect calls (e.g., `jmp eax`)
  - Course-grained: static call graphs
  - Finer-grained: type-based, Input-based, etc.
- Backward CFI: protecting returns (e.g., `ret`)
  - Safe/shadow stack

→ CPU overheads: 5-20% in forward CFIs, 1-5% in backwards CFIs

# Trends of Vulnerability Classes

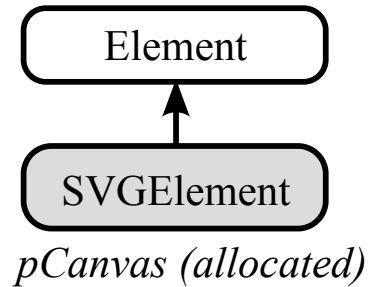


Ref. [Exploitation Trends: From Potential Risk to Actual Risk, RSA 2015](#)

# Three Emerging Classes of Vulnerabilities


1. Type confusion: e.g., bad casting in C++
2. Race condition: e.g., double fetching
3. Uninitialized read: e.g., struct padding issues

# Type Casting in C++



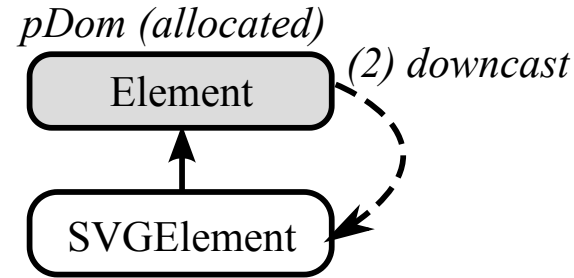
```
1 | class SVGElement: public Element { ... };  
2 |  
3 | SVGElement *pCanvas = new SVGElement();
```

# Upcasting and Downcasting in C++



```
1 // (1) valid upcast from pCanvas to pElem
2 Element *pElem = static_cast<Element*>(pCanvas);
3
4 // (2) valid downcast from pElem to pCanvasAgain (== pCanvas)
5 SVGElement *pCanvasAgain = static_cast<SVGElement*>(pElem);
```

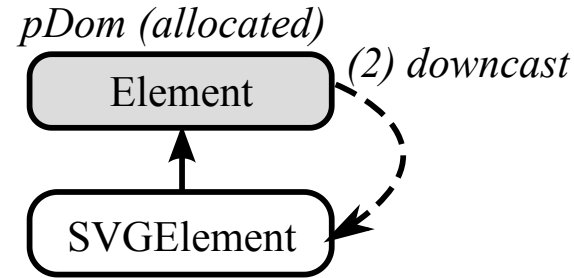
# Dynamic Casting in C++ (Downcasting)



```
1 | Element *pDom = new Element();  
2 |  
3 | // (3) invalid downcast with dynamic_cast, but no corruption  
4 | SVGElement *p = dynamic_cast<SVGElement*>(pDom);  
5 | if (p) {  
6 |     p->m_className = "my-canvas";  
7 | }
```



# Static Casting in C++ (Downcasting)

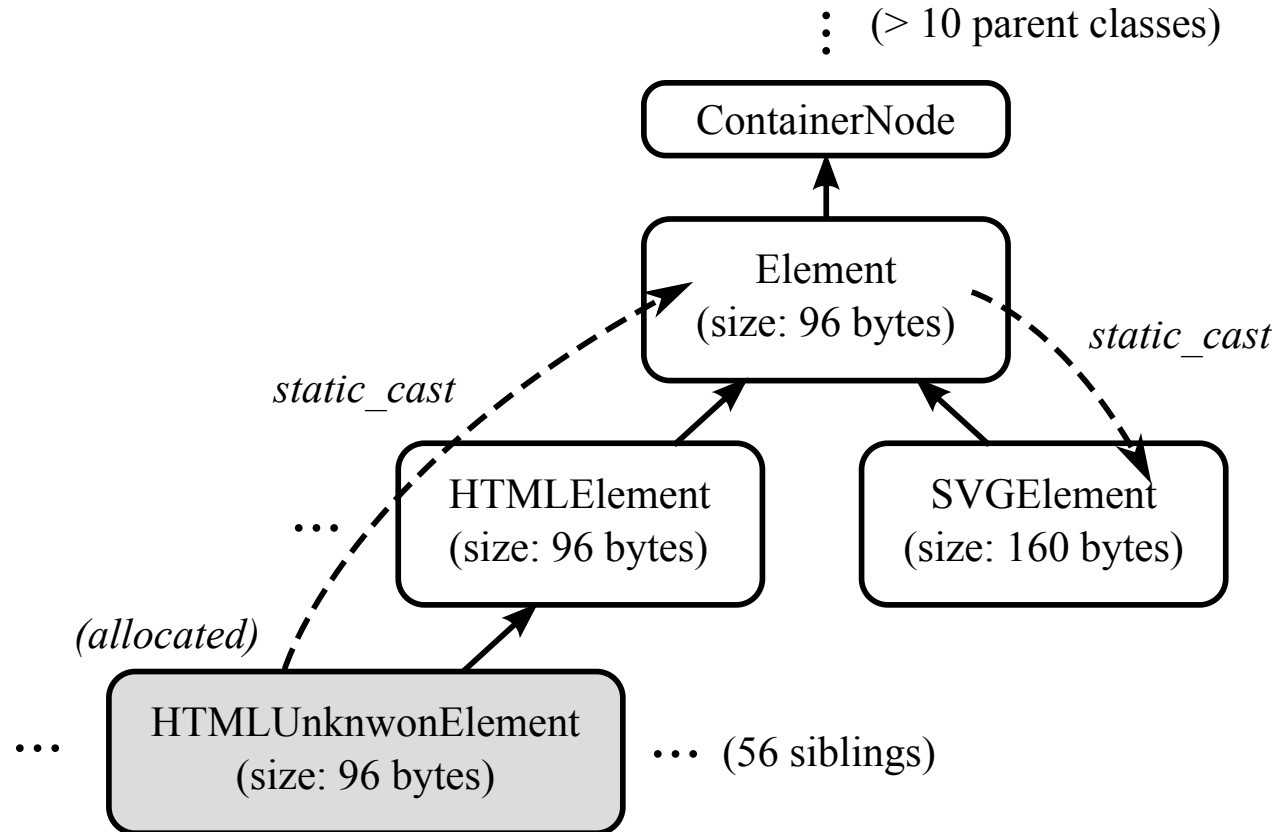


```
1 // (4) BUG: invalid downcast
2 SVGElement *p = static_cast<SVGElement*>(pDom);
3
4 // (5) leads to memory corruption!
5 p->m_className = "my-canvas";
```

# Type Confusion: Bad Casting in C++

- Two (important) type castings:
  1. `static_cast<>`: verify the correctness at compilation
  2. `dynamic_cast<>`: verify the correctness at runtime
- Incorrect downcasting (via `static_cast<>`) results in memory violation
- `dynamic_cast<>` requires RTTI, incurring non-trivial perf. overheads

# CVE-2013-0912 and Pwn2Own 2013



# Exercise: Real-world Examples

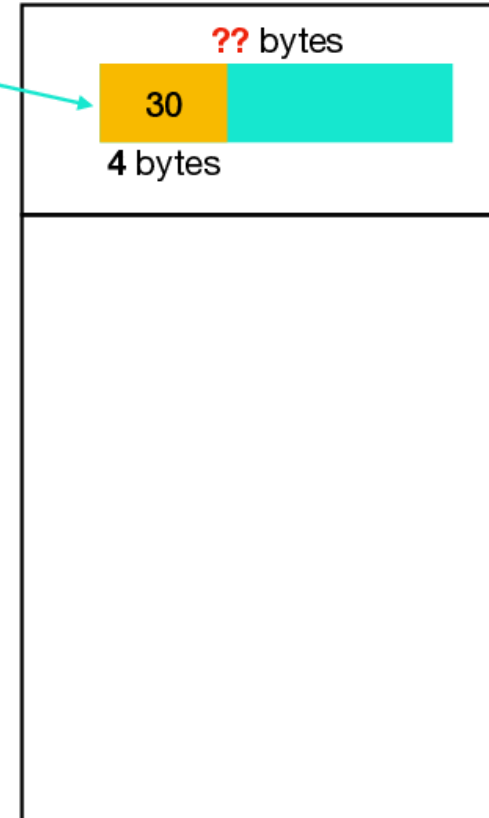
- Ex1. Linux Perf (CVE-2009-3234/+)
  - → Race condition: double fetching in Linux
- Ex2. Linux USB (CVE-2016-4482)
  - → Uninitialized read via struct padding

# Double Fetching Vulnerabilities

- A special form of race conditions: user vs. kernel spaces
- Leading to information leaks, buffer overflows, etc.

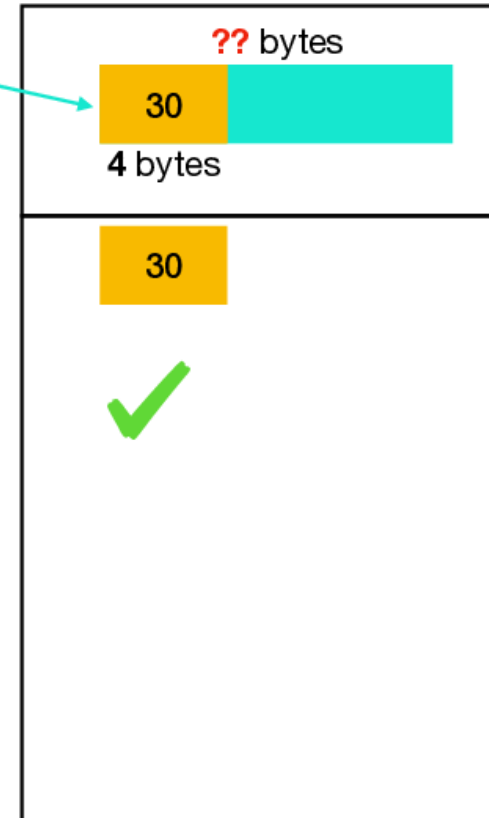
# Example: perf\_copy\_attr() Explained

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr __user *uattr,
3    struct perf_event_attr *attr) {
4
5   u32 size;
```



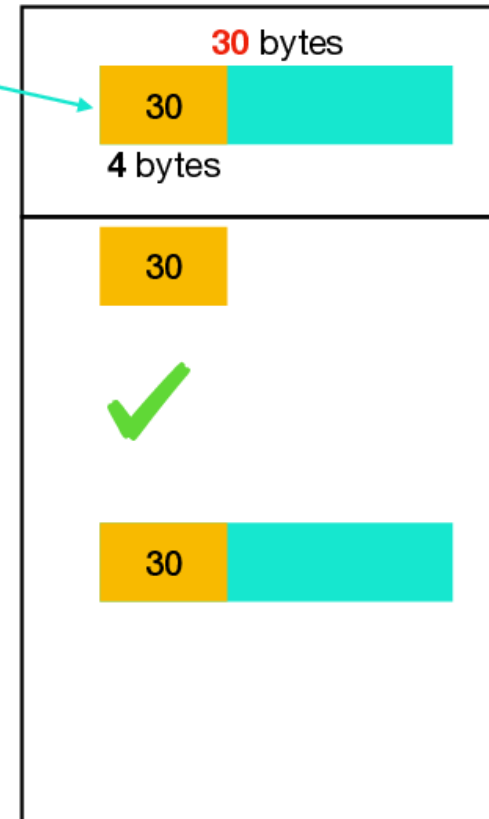
# Example: perf\_copy\_attr() Explained

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr __user *uattr,
3    struct perf_event_attr *attr) {
4
5     u32 size;
6
7     // first fetch
8     if (get_user(size, &uattr->size))
9         return -EFAULT;
10
11    // sanity checks
12    if (size > PAGE_SIZE ||
13        size < PERF_ATTR_SIZE_VER0)
14        return -EINVAL;
```



# Example: perf\_copy\_attr() Explained

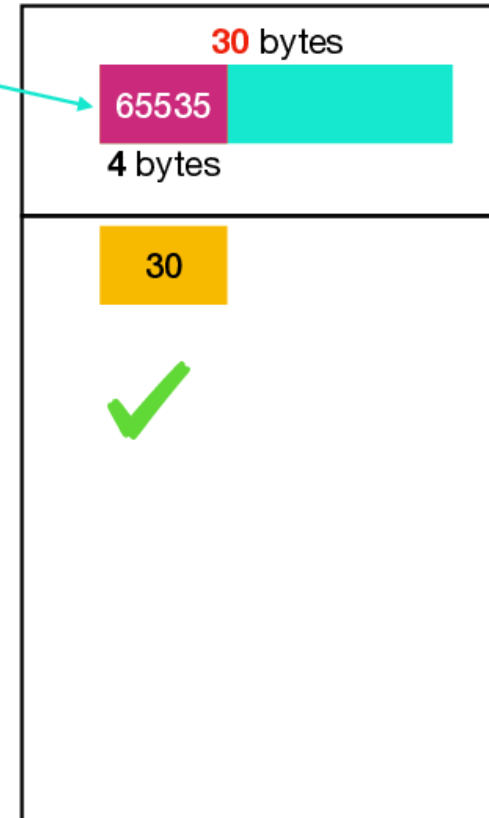
```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr __user *uattr,
3    struct perf_event_attr *attr) {
4
5     u32 size;
6
7     // first fetch
8     if (get_user(size, &uattr->size))
9         return -EFAULT;
10
11    // sanity checks
12    if (size > PAGE_SIZE ||
13        size < PERF_ATTR_SIZE_VER0)
14        return -EINVAL;
15
16    // second fetch
17    if (copy_from_user(attr, uattr, size))
18        return -EFAULT;
19
20    .....
21 }
```






# BUG: Racing in Userspace

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr __user *uattr,
3    struct perf_event_attr *attr) {
4
5     u32 size;
6
7     // first fetch
8     if (get_user(size, &uattr->size))
9         return -EFAULT;
10
11     // sanity checks
12     if (size > PAGE_SIZE ||
13         size < PERF_ATTR_SIZE_VER0)
14         return -EINVAL;
```



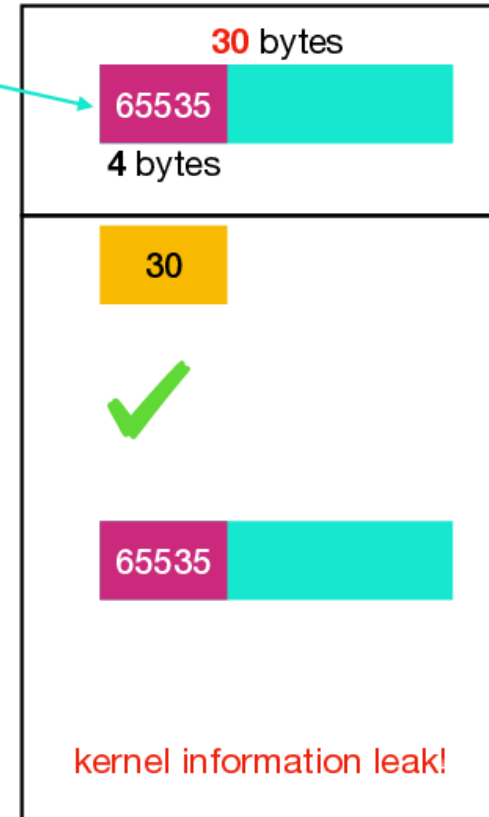
# BUG: “Double” Fetching from Kernel

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr __user *uattr,
3    struct perf_event_attr *attr) {
4
5     u32 size;
6
7     // first fetch
8     if (get_user(size, &uattr->size))
9         return -EFAULT;
10
11    // sanity checks
12    if (size > PAGE_SIZE ||
13        size < PERF_ATTR_SIZE_VER0)
14        return -EINVAL;
15
16    // second fetch
17    if (copy_from_user(attr, uattr, size))
18        return -EFAULT;
19
20    .....
21 }
```



# BUG: Trigger Incorrect Memory Copy

```
1 static int perf_copy_attr_simplified
2   (struct perf_event_attr __user *uattr,
3    struct perf_event_attr *attr) {
4
5     u32 size;
6
7     // first fetch
8     if (get_user(size, &uattr->size))
9         return -EFAULT;
10
11    // sanity checks
12    if (size > PAGE_SIZE ||
13        size < PERF_ATTR_SIZE_VER0)
14        return -EINVAL;
15
16    // second fetch
17    if (copy_from_user(attr, uattr, size))
18        return -EFAULT;
19
20    .....
21 }
22
23 // BUG: when attr->size is used later
24 copy_to_user(ubuf, attr, attr->size);
```



# Fixing Double Fetches

1. Partially reading after the size attribute
2. Ensuring the atomicity of previous checked size

```
1  // @f12f42acdbb577a12eecfcebbsbec41c81505c4dc
2  ret = get_user(size, &uattr->size);
3  ret = copy_from_user(attr, uattr, size);
4  ...
5
6  // overwrite with the sanity-checked size
7  + attr->size = size;
```

# CVE-2009-3234: Buffer Overflow!

```
1  /* If we're handed a bigger struct than we know of,  
2   * ensure all the unknown bits are 0.  */  
3  if (size > sizeof(*attr)) {  
4      ...  
5      for (; addr < end; addr += sizeof(unsigned long)) {  
6          ret = get_user(val, addr);  
7          if (ret)  
8              return ret;  
9          if (val)  
10             goto err_size;  
11     }  
12 }  
13 ret = copy_from_user(attr, uattr, size); // Q. size?
```

# CVE-2009-3234: Buffer Overflow!

```
1  /* If we're handed a bigger struct than we know of,  
2   * ensure all the unknown bits are 0.  */  
3  if (size > sizeof(*attr)) {  
4      ...  
5      for (; addr < end; addr += sizeof(unsigned long)) {  
6          ret = get_user(val, addr);  
7          if (ret)  
8              return ret;  
9          if (val)  
10             goto err_size;  
11     }  
12 +   size = sizeof(*attr);  
13     }  
14     ret = copy_from_user(attr, uattr, size); // Q. size?
```

# Exercise: Real-world Examples

- Ex1. Linux Perf (CVE-2009-3234/+)
  - → Race condition: double fetching in Linux
- Ex2. Linux USB (CVE-2016-4482)
  - → Uninitialized read via struct padding

# CVE-2016-4482: Linux USB

```
1 static int proc_connectinfo(struct usb_dev_state *ps,
2                             void __user *arg) {
3     struct usbdevfs_connectinfo ci = {
4         .devnum = ps->dev->devnum,
5         .slow = ps->dev->speed == USB_SPEED_LOW
6     };
7
8     if (copy_to_user(arg, &ci, sizeof(ci)))
9         return -EFAULT;
10    return 0;
11 }
```



# Padding Issues in Struct

```
struct usbdevfs_connectinfo {  
    unsigned int devnum;    // 4 bytes  
    unsigned char slw;      // 1 bytes  
};
```

```
sizeof(struct usbdevfs_connectinfo) == 8
```

```
|<----- 8B ----->|  
[devnum      ][slw][padding]  
|<-- 4B  -->|<1B>|<--3B-->|
```

Ref. [Proactive Kernel Memory Initialization to Eliminate Data Leakages](#)

# Struct Padding: No Proper Way to Initialize



§6.2.6.1/6 (C11, ISO/IEC 9899:201x)

*When a value is stored in an object of structure (...), the bytes of the object representation that correspond to any padding values.*

```
struct usbdevfs_connectinfo ci = {  
    .devnum = ps->dev->devnum,  
    .slow = ps->dev->speed == USB_SPEED_LOW  
};
```

# CVE-2016-4482: Patch

- Zero-ing out via memset(): initializing two times!

```
1  static int proc_connectinfo(struct usb_dev_state *ps,  
2                               void __user *arg) {  
3  +   struct usbdevfs_connectinfo ci;  
4  +  
5  +   memset(&ci, 0, sizeof(ci));  
6  +   ci.devnum = ps->dev->devnum;  
7  +   ci.slow = ps->dev->speed == USB_SPEED_LOW;  
8  +   ...  
9  }
```

# Summary

- Modern attacks are using **ROP** to defeat DEP/NX
- Learn three emerging, critical classes of vulnerabilities
  - **Type confusion**: e.g., bad casting in C++
  - **Race condition**: e.g., double fetching
  - **Uninitialized read**: e.g., struct padding issues

# References

- [CVE-2016-4482](#)
- [CVE-2013-0912](#)
- [CVE-2009-3234](#)
- [Proactive Kernel Memory Initialization to Eliminate Data Leakages](#)
- [Precise and Scalable Detection of Double-Fetch Bugs in OS Kernels](#)
- [Return-oriented Programming](#)
- Unisan: [paper/slides](#)